

Title

COLDSTART: Anticipating health-system crises before the outcome is observable

1. Summary

When a health crisis begins, the quantity decision-makers need to forecast (presentations, demand, occupancy) is precisely the one for which almost no local outcome data yet exist. Context variables abound, but the outcome series is initially a handful of noisy observations, so systems that perform well on years of local history are weakest when early decisions matter most.

There is another source of quantitative information: published evidence from analogous events elsewhere, thousands of studies reporting transmission parameters, weather-demand associations and surge magnitudes. It is rarely tested as formal prior information, because of the two steps that matter most: **reliable extraction of quantitative estimates, and their transport to the new setting**. A confidently wrong prior is worse than no prior, because it is most influential when local data cannot yet correct it.

This project asks the central question directly: **can published quantitative evidence provide useful information when local outcome data are insufficient at crisis onset, and can we detect early when it should not be trusted?** **LiteRev-Evidence**, the prototype I built, shows the pipeline, from query across the open literature through screening and provenance-tracked extraction to pooled parameter distributions. The project turns the prototype into a validated instrument and tests it: I will (1) establish the error structure of automated extraction and propagate it into evidence-derived priors; (2) develop an uncertainty-aware latent-regime framework for health-system escalation, with those priors and a short-window resilience signal; (3) test, by strict rolling-origin reconstruction of each historical moment's information, whether the priors improve probabilistic forecast skill during the cold-start phase, how fast that value decays, and when they cause harm; and (4) determine whether any gain is large enough to change operational decisions. The decisive experiment is a pre-specified ladder of models, from local baselines to literature-derived priors and adaptive borrowing that discounts them under prior-data conflict; one rung-to-rung contrast is confirmatory, repeated once, sequentially, as the heat generalisation test.

Validation uses contrasting crisis archetypes: respiratory epidemics carry the confirmatory test, heatwaves the sequential generalisation test, Geneva legionellosis a scoped year-4 extension. The outcome is intentionally falsifiable: if literature-derived priors improve cold-start forecasts, the project establishes when and by how much; if not, or if their harm cannot be detected early enough, it establishes a boundary condition for automated evidence-informed forecasting. Either way, the deliverable is a validated answer and an open framework, not another untested forecasting platform. The project rests on an unusual combination: fifteen years in quantitative finance on rare-event risk, regime models and stress testing; **LiteRev-Evidence**; Swiss hospital and emergency data. Ambizione turns these into an independent research programme, led by me, at the intersection of evidence synthesis, quantitative modelling and emergency public health.

2. Proposed research

2.1 Current state of research in the field

2.1.1 Forecasting health-system crises: strong in steady state, weak at onset

Syndromic surveillance detects departures from expected baselines [Farrington 1996; Noufaily 2013], and forecast ensembles perform strongly in data-rich settings [Cramer 2022; Sherratt 2023]. Dispatch records lead confirmed surveillance: the 2009 A(H1N1) autumn onset was identified eight days in advance across three European regions [Rosenkötter 2013], long call series track influenza-like illness [EMS-ILI 2025], and our systematic review maps the field [Edjinedja 2026].

The decisive limitation is **history dependence**: data-adaptive models need enough local observations to learn seasonality, weather response and crisis dynamics. What to do in the first weeks of a novel or displaced crisis, when the local outcome series is short and unstable, remains open.

The cold-start problem is specifically a labelled-outcome problem. At crisis onset, context variables abound but the outcome to be forecast does not, and abundance itself can mislead: Google Flu Trends overestimated influenza activity by more than a factor of two [Lazer 2014]. Early signals do not close the gap: wastewater needs a shedding-to-incidence relationship that itself comes from external evidence; transfer learning needs contemporaneous comparable observations; mechanistic models need hand-set parameters from a few studies. The common issue is **how to use external quantitative information without pretending it is perfectly transferable**.

2.1.2 The unused resource: published evidence as quantitative prior information

Thousands of studies report weather-demand associations, surge multipliers, transmission parameters and length-of-stay distributions, and Bayesian borrowing is well developed: power priors [Ibrahim 2000], commensurate priors [Hobbs 2011], meta-analytic-predictive priors with robust conflict protection [Schmidli 2014]. Informative priors have been used to forecast disease under sparse local data [Cook 2023]. What has not been established is narrower and harder: whether an **automatically constructed** evidence prior, with extraction error and transportability uncertainty propagated into it, improves **operational cold-start forecasts under strict historical information constraints**, and whether harmful borrowing is detected early enough to act.

Extraction. Automated extraction makes large-scale synthesis feasible, but numerical extraction remains less reliable than categorical (reported accuracy roughly 47–88% versus 74–96%), and omissions dominate errors [Shankar 2026]. The question is whether extraction and pooling preserve the dispersion a calibrated prior needs.

Transportability. Even perfectly extracted estimates may not transfer across populations, case definitions, health systems or policy regimes. Formal transportability tools exist [Bareinboim 2016; Dahabreh 2019; Degtiar 2023] but are not integrated into any operational cold-start framework.

2.1.3 Representing escalation as a state, not only a point forecast

Emergency operations are interested in state (routine, elevated, strained, critical) rather than only tomorrow's count. Latent-state representations are established in surveillance: Poisson hidden Markov models distinguish epidemic from non-epidemic periods [Le Strat 1999; Watkins 2009]; extreme-value methods have been applied to Swiss hospital visits [Coles 2001; Ranjbar 2022]. Here

they serve one purpose: a latent ordinal state as the **common representation for testing evidence borrowing at crisis onset**, with an extreme-value component for rare critical exceedances. Critical-slowness theory (rising variance and lag-1 autocorrelation before some transitions [Scheffer 2009]; epidemic applications [O'Regan 2013; Brett 2018; Southall 2021]) uses the shape of a short recent series: a **secondary information channel** tested empirically, not a universal early-warning mechanism.

2.1.4 From forecast accuracy to decision value

Proper scoring rules reward calibration [Gneiting 2007], but a calibrated forecast is irrelevant if it changes no decision; decision-analytic methods evaluate predictions under explicit consequences [Vickers 2006]. 72–99% of reported clinical alarms have been false or non-actionable [Winters 2018]; mortality-calibrated heat thresholds need not match demand thresholds [Lee 2021]; pollution-asthma emergency associations lack recent Swiss acute-episode studies [Zheng 2015]. Methods are therefore evaluated first statistically, then under an elicited operational loss structure.

2.1.5 The specific gap addressed by this project

Ongoing programmes occupy the neighbouring ground: multi-model forecasting hubs run routinely (ECDC RespiCast, US CDC FluSight), Horizon Europe funds epidemic-intelligence infrastructure (including GeoAI4EI, in which I take part), and the Franco-Swiss GESICA programme builds cross-border crisis intelligence. These programmes develop surveillance, forecasting and decision-support capability from local data; their stated objectives do not include testing whether automatically synthesised external evidence should enter a forecast as formal prior information at local cold start. The four pieces exist separately. What has not been established is whether they connect around the **cold-start question**: whether systematically extracted evidence improves probabilistic forecasting before local outcome data become informative, and whether harmful borrowing is detectable early enough to discount.

2.2 Current state of personal research and competences required for the project

My route is unusual and is why the project is tractable: fifteen years in quantitative finance (risk modelling, extreme-value estimation for non-Gaussian tails, regime and factor models, stress testing), then a doctorate in biomedical sciences at Geneva (defended 18 December 2023).

2.2.1 Prediction under sparse and imperfect information

In **Orel et al., *PLoS ONE* 2022** I predicted individual HIV status from socio-behavioural characteristics across East and Southern Africa, establishing where models transported between countries. **Merzouki et al., *PeerJ* 2021** and **Esra et al., *JAIDS* 2023** developed the same theme. I now **senior-author** that line (last author, **Ng'ambi et al., *accepted***: machine-learning classification of cardiovascular disease history across harmonised WHO STEPS surveys). That is this proposal's transportability problem, met first elsewhere and moved from conducting to directing.

2.2.2 Automated evidence synthesis

Orel et al., *J Med Internet Res* 2023 introduced **LiteRev**, an automated literature-review tool: natural-language or Boolean query across the open-access databases, deduplication, corpus mapping, iterative relevance suggestion. I led its development with **Aziza Merzouki** (PhD, computer science) and secured development funding on my own initiative: **CHF 30,000** (UNIGE), **CHF 10,000** (Venture Kick), **CHF 20,000** (Mimosa), outside any group grant. The AI-in-EMS systematic

review (Edjinedja, Larribau, Orel et al., submitted 2026), within GESICA, used it to structure 138 retained publications.

2.2.3 Outbreak and health-system modelling in Switzerland

Orel et al., *CMI Communications* 2024 compared clinical severity between Delta and Omicron sub-lineages in a Swiss tertiary centre. Estill et al., *F1000Research* 2020 built age-structured scenario models for the Swiss SARS-CoV-2 epidemic under planning time pressure: the experience this proposal's question comes from. I contributed to WHO African Region reporting and seroprevalence estimation (*Nat Commun*, Nwosu et al., 2021).

2.2.4 The prototype and test bed: LiteRev-Evidence

Since 2024 I have developed **LiteRev-Evidence**, extending LiteRev from retrieval and screening into structured quantitative extraction and modelling: the full review pipeline in working form, from a query searched live across the open literature, through deduplication, PRISMA-accounted screening and PICO extraction, to **provenance-tracked, quality-scored extraction pooled into parameter distributions** and propagated through ensemble simulation, with per-scenario **living reviews re-run daily**: the literature-to-prior mechanism this proposal interrogates, in prototype form. Operational scenarios are elaborated with emergency-medicine partners.

For GESICA I built the Geneva-Vaud-Neuchâtel data foundation: **77 notifiable diseases in eight model classes** by transmission mode, and a referenced inventory of **28 surveillance sources (23 infectious, 5 environmental/non-infectious)** each documented for institution, coverage, resolution, latency, access and quality. It is why the validation domains are chosen by **model class**, and why this proposal rests on a mapped data landscape rather than an assumed one.

The prototype makes the research feasible rather than aspirational; the validated instrument it becomes is itself a deliverable.

2.2.5 Linked data on a contrasting crisis archetype

I lead the data work on a cantonal study **already under way** (BASEC 2026-00324, ethics granted) linking confirmed legionellosis cases in Geneva to individual domestic hot-water installations: a linkage to my knowledge unique, supplying a waterborne archetype whose dynamics differ fundamentally from a respiratory epidemic's.

Competences. Bayesian hierarchical and regime-switching estimation, extreme-value modelling, stress testing under misspecification, forecast evaluation and decision analysis come from fifteen years of quantitative risk work and form the core of WP2-WP3; extraction, retrieval and quality-weighted pooling from LiteRev-Evidence, which I built; domain knowledge from the doctorate, the Swiss COVID-19 work and GESICA.

My other commitments and their delimitation are in §2.6; what I have not yet had is a programme of my own with time to run it.

2.3 Detailed research plan

2.3.1 Objectives and hypotheses




Figure 1: cold-

start problem, evidence-borrowing hypothesis, evaluation ladder.

To determine **whether, and under what conditions, published quantitative evidence provides useful information when local outcome data are insufficient at the onset of a health-system crisis, and whether the resulting forecasts change decisions.**

One central hypothesis, H3a; everything else is subordinate to it:

	Role	Statement
H1	Validation	Can the evidence be trusted enough to use?
C2	Adequacy criterion	Is the state representation fit to compare borrowing in?
H3a	Central hypothesis	Do evidence-derived priors improve cold-start forecast skill?
H3b	Robustness	Is adaptive borrowing safe when the prior is wrong?
H3c	Secondary channel	Do resilience indicators add information beyond the prior?
H4	Decision value	Is the gain large enough to change an operational choice?

O1: Make published evidence usable without hiding its uncertainty

H1. Automated extraction will systematically **understate the dispersion of the evidence-derived distribution** (reported within-study uncertainty and between-study heterogeneity), producing priors that are too concentrated; an explicit measurement-error layer will recover enough of the missing dispersion to construct usable priors.

The direction is mechanistic: omissions dominate extraction errors [Shankar 2026] and shrink estimated dispersion systematically; if the overconfidence cannot be corrected, the project establishes a boundary condition.

O2: Represent escalation in a form that separates state from the point forecast

C2, model adequacy criterion. The latent-state representation must yield **identifiable** parameters and **calibrated** escalation-state probabilities at matched false-alarm rates. Its role is the common state representation in which borrowing strategies are compared, not a claim that regime switching generally beats thresholding a point forecast.

C2 is verified rather than discovered: T2.1's identifiability study and T3.3's calibration checks either establish adequacy or trigger the pre-specified ordinal fallback; either outcome leaves H3a intact.

Extreme-value modelling represents the critical tail; critical-slowness indicators are **supporting covariates** tested against level and trend.

O3: Test the cold-start hypothesis and map failure

H3a. Evidence-derived priors improve probabilistic forecast skill during the early phase of a crisis, with the advantage declining as local observations accumulate.

H3b. Adaptive borrowing that discounts the evidence when prior-data conflict emerges is **non-inferior** to fixed borrowing under well-specified priors, within a pre-specified margin Δ on the CRPS skill score, and is **superior** to fixed borrowing under deliberately misspecified priors.

H3b is two-sided: non-inferiority where the evidence is sound, superiority where it is not; Δ is fixed at the second registration point, justified against the rung 3 $\rightarrow$ rung 4 effect the study is powered to detect.

H3c. Resilience indicators add predictive information beyond the evidence-derived prior and the local level/trend signal when the outcome history is short.

Primary confirmatory comparison (one, stated once): rung 4 (fixed evidence-derived priors) against rung 3 (weakly informative priors), by **CRPS skill score**, over the pre-specified cold-start window, **on respiratory episodes only** (procedure and registration: T3.3). Heat repeats the identical contrast as a **sequential generalisation test**, run only if the respiratory test is met; the fixed order controls the family-wise error rate and follows the science: respiratory evidence is richest, heat transport hardest. A respiratory-positive, heat-negative result is a boundary condition on transportability. The **shape of the advantage over elapsed local data** is reported: it should decay to nothing.

O4: Establish whether predictive improvement is decision-relevant

H4. Decision-analytic evaluation under the losses and escalation thresholds of emergency responders can rank modelling strategies differently from generic accuracy criteria; the evidence-derived strategy is useful only when its gain crosses a decision threshold.

These are downstream tests of value; shadow mode is an extension, not a prerequisite.

Validation domains

The domains span **two contrasting model classes** from my GESICA classification. Both arms score **the same quantity, built the same way**: daily emergency demand from the CASU-144 series, restricted to the cause classes the archetype's evidence base concerns (T3.0): respiratory-related demand for epidemics; **heat-sensitive demand** for heat, since Swiss evidence places heat effects in dehydration, renal and psychiatric admissions, with a weak respiratory effect at older ages [Schulte 2024; Ragettli 2019]. Matching each arm's outcome to its evidence base makes the generalisation test a test of borrowing, not of the outcome definition.

Archetype	Role	Dynamics	Outcome and data
Respiratory epidemic	Primary confirmatory	Transmissible, multi-wave, seasonal	Respiratory-related demand; COVID-19, influenza, RSV
Heatwave	Sequential generalisation	Environmental, short, sharply peaked	Heat-sensitive demand, same construction; MeteoSwiss exposures
Waterborne outbreak	Year-4 extension	Common-source, environmental	Geneva legionellosis linked to installations

What the project does not claim

It does **not** aim to outperform forecast hubs in the data-rich regime, assume literature-derived priors are beneficial, claim critical slowing down as universal, or promise a deployed clinical alarm system by month 48. The contribution is narrower: **whether accumulated quantitative evidence can earn a formal role in forecasting before local outcome data become informative, and a map of when it should not be trusted.**

2.3.2 Work packages and methods

WP1 establishes whether the evidence can be trusted; WP2 supplies the common state representation; WP3 is the decisive cold-start evaluation; WP4 translates predictive differences into operational value.

WP1: From published evidence to usable priors (M1–M20)

T1.1: Define the evidence target and register the protocol (M1–M4)

Core (benchmarked in full): the quantities the demand model consumes: **weather-demand associations, surge magnitudes** (peak-to-baseline ratios), **length-of-stay / occupancy distributions**. **Secondary (only if core completes on schedule):** transmission parameters and shedding-to-incidence conversions; neither is on the critical path. Inclusion criteria, effect measures, uncertainty representation and transportability variables are fixed first.

T1.2: Build the quantitative extraction benchmark (M3–M9)

Two independent expert extractors extract target quantities from a stratified random sample of **300** publications, with adjudication; stratification covers parameter class, reporting quality, study design and **source type** (journal article versus timestamped situational report), so error is characterised separately for the faster sources of T3.2's secondary prior variant. The benchmark is an open deliverable.

T1.3: Characterise automated extraction error (M7–M14)

Compare automated extraction with the benchmark on point estimates, uncertainty, omissions and between-study dispersion, with sensitivity to the underlying model version.

T1.4: Construct uncertainty-aware evidence distributions (M12–M18)

Represent extraction error explicitly as measurement error; compare inverse-variance and quality-weighted pooling, meta-analytic-predictive priors and power-prior discounting. Output: a prior with provenance, extraction uncertainty and a documented borrowing weight.

T1.5: Transportability screen (M15–M20)

Characterise effect modifiers and study-setting differences relevant to Geneva; where transport is weak, widen or discount the prior, producing the metadata WP3's conflict analysis needs.

Deliverables. D1.1 open benchmark; D1.2 error analysis; D1.3 evidence-to-prior library.

Risk. Manual extraction is expensive; if expert time binds, the benchmark narrows to the core parameter classes.

WP2: A parsimonious model of health-system escalation (M1-M28)

T2.1: Specify and identify the state model (M1-M9)

A Bayesian hierarchical Markov regime-switching model [Hamilton 1989] with an ordinal latent state $S(t) \in \{\text{routine, elevated, strained, critical}\}$ observed through emergency calls, ED presentations and ICU occupancy; pre-specified weather, calendar and epidemic covariates; series-specific observation models sharing the state. Before real-data fitting: simulation-based identifiability and recovery; ordering constraints resolve label switching; if separation is insufficient, the fallback is an ordinal state-space formulation; the criterion is recovery of states and transitions, not visual fit.

T2.2: Represent the critical tail (M6-M14)

Peaks-over-threshold/generalised Pareto modelling for rare exceedances, coupled to the critical-state probability, threshold sensitivity reported: a supporting representation, not a separate objective.

T2.3: Introduce evidence-derived priors (M10-M20)

Map WP1 distributions to the parameters where published evidence is relevant (weather effects, surge magnitudes, transition/recovery characteristics); compare weakly informative, fixed evidence-derived and adaptive robust borrowing, distinguishing the value of evidence from that of the regime representation. The adaptive specification is a robust mixture of an evidence-derived and a weakly informative component, prior-data conflict recorded explicitly. Power-prior and commensurate-prior approaches are sensitivity comparators.

T2.4: Add resilience indicators as a secondary information channel (M12-M20)

Rolling variance and lag-1 autocorrelation, entered as optional covariates on transition dynamics: do they add information beyond level/trend and the evidence prior? A null result is interpretable.

T2.5: Calibration and implementation (M20-M28)

Calibration methods suited to temporal dependence; a conformal component [Angelopoulos 2023; Barber 2023] as a robustness layer if simulation confirms its assumptions, not a headline claim. Release a documented reference implementation integrated with LiteRev-Evidence.

Simulation is prior information, not data. Fitting to literature-parameterised trajectories as if independent observations would count the same prior information twice and disable the prior-data discrepancy diagnostic H3b depends on. Mechanistic simulation serves only identifiability studies, structural constraints and prior predictive checking; generated trajectories never tighten the evidence prior.

Deliverables. D2.1 model specification and identifiability study; D2.2 open implementation; D2.3 methodological paper.

Risk. Weak regime separation. *Mitigation:* simulation first; the ordinal fallback; report the identifiability boundary as a result rather than tuning until success.

WP3: The decisive cold-start experiment (M12-M42)

T3.0: Outcome hierarchy, data-access gate and episode eligibility (M1-M14)

Fixed before any evaluation is designed, never revisited in response to performance:

Outcome. The primary outcome is **daily respiratory-related emergency demand derived from CASU-144 records**, not a raw call count: "respiratory-related" is a pre-registered classification of

recorded call reasons and urgency levels, with sensitivity to the construction reported. **The heat arm scores heat-sensitive demand:** the same series and construction, restricted to cause classes fixed at registration (dehydration, renal, psychiatric, per the Swiss evidence), with respiratory-restricted and 75+ sensitivities. **The primary claim is scoped to what the series measures:** the operational demand signal observed by the dispatch system. ED presentations and ICU occupancy enter as **additional observation channels on the shared latent state**; wastewater, sentinel consultations and weather as covariates; the hierarchy governs only which series H3a is scored on. **Data-access gate.** Criteria fixed in advance for **historical depth, temporal resolution, reporting latency and completeness**; the primary outcome is selected at a pre-specified checkpoint (month 12), on those criteria alone.

Episode eligibility. An episode enters the confirmatory evaluation only if all of these are prospectively reconstructable: (1) a detectable onset under the prospective onset rule; (2) enough pre-onset history for its rolling baseline; (3) enough post-onset observations at the primary horizon; (4) the external evidence **as it stood at the historical origin**; (5) no leakage of future information; (6) separation from adjacent episodes. Failing episodes remain available for sensitivity analysis only. **At demand level, co-circulating pathogens form one episode:** a winter with concurrent influenza and RSV is one surge, not two.

Two registration points. The hierarchy, gate criteria and eligibility rule are registered now. The window N , the archetype-specific horizons and the margin Δ are registered after the checkpoint and the episode inventory, but before any evaluation runs, by the pre-declared selection rules of T3.3. Earlier would be guesswork; after any look at performance, indefensible.

T3.1: Assemble the retrospective information set ($M12-M20$)

Harmonise the CASU-144 series with the additional channels and covariates; quantify completeness and delay; model right truncation/nowcasting so incomplete recent reporting is not mistaken for falling demand [Höhle 2014; McGough 2020]. Degraded reporting under strain could itself mimic an early-warning signal, so missingness and delay are characterised explicitly.

T3.2: Reconstruct the true information set ($M18-M30$)

For each historical onset, create successive forecast origins using **only information available at that date**. **Admissibility as a prior is defined by referent, not venue:** prior inputs are statements about *other* populations, places or past events. **The confirmatory prior uses peer-reviewed literature and preprints only**, under the rolling cut-off. Timestamped situational reporting on the ongoing event *elsewhere* (WHO Disease Outbreak News, ECDC rapid risk assessments) enters as a **pre-specified secondary prior variant**: it passes through the same extraction and measurement-error pipeline but stays out of the confirmatory contrast: H3a tests literature borrowing, not dynamic evidence fusion. **Text describing the local event is excluded from any prior:** it is a noisy measurement of the outcome H3a is scored against. Every input carries an index timestamp enforcing the rolling cut-off. Forecast horizons 7, 14 and 28 days, the archetype-specific primary horizon fixed at the second registration point. The cold-start window is defined by elapsed local outcome observations **after a pre-defined real-time onset criterion**; the onset rule may use only variables available at the origin, never any future information.

T3.3: Evaluate a pre-specified model ladder (M20-M34)

At each origin compare: (1) seasonal/naive local baseline; (2) established short-baseline surveillance method; (3) regime model, weakly informative priors; (4) the same model, fixed evidence-derived priors; (5) adaptive borrowing with conflict monitoring; (6) adaptive borrowing plus resilience indicators.

Confirmatory testing procedure (fixed-sequence).

1. **Test 1 (primary).** Rung 4 vs rung 3, CRPS skill score, **respiratory** episodes, cold-start window N , at $\alpha = 0.05$ two-sided. Passing requires paired-permutation $p < \alpha$ *and* a lower confidence bound above the pre-registered minimal relevant improvement.
2. **Test 2 (generalisation).** The identical contrast on **heat** episodes, scored on heat-sensitive demand (T3.0), at the same α , **conducted only if Test 1 passes**.
3. **If Test 1 fails**, H3a is not supported, Test 2 is exploratory, and the project's result is the failure map and the boundary condition.

All other ladder contrasts are secondary. **Rung 3 is pinned in the registration**, with a pre-declared set of vaguer and tighter alternatives reported as a sensitivity band: the comparator cannot be tuned into a straw man after the fact.

Primary endpoint: the CRPS skill score of rung 4 versus rung 3 over the cold-start window. The unit of inference is the **episode** (origins are repeated measures within it); eligible episodes are few (13–14 respiratory candidates before screening), so inference is by **paired permutation over episodes**, block bootstrap alongside, both pre-specified, disagreement reported. **The design is powered by simulation-based operating characteristics at episode level** (type I error, power, interval width under realistic within-episode dependence and heterogeneity), run before the second registration point. The same simulation fixes, in advance, the **minimum number of eligible episodes** below which no confirmatory claim is made and the primary analysis becomes estimation, and the **selection rules** for N and Δ : N the smallest grid value meeting the operating-characteristic criterion, Δ set by a pre-declared rule on the simulated effect distribution, both chosen by algorithms blind to forecast performance. Secondary endpoints: log score, calibration (PIT, coverage), escalation detection at matched false-alarm rates. **H3b's Δ is fixed here, before any evaluation:** non-inferiority is a CRPS deficit no greater than Δ ; superiority is tested on T3.4's misspecified priors.

T3.4: Map benefit and failure (M28-M38)

Identify episodes where borrowing improves or worsens forecasts; characterise failure by population, outcome-definition, system, policy and temporal mismatch, extraction uncertainty and prior-data conflict, with deliberately misspecified priors as a stress test. The safety question: is harmful borrowing detectable early enough for adaptive discounting to limit its impact? The output is a **failure map**, not an average.

T3.5: Test generalisation (M34-M42)

The sequential heat test and, resources permitting, the legionellosis extension (year 4, not required for the central conclusion).

Deliverables. D3.1 reproducible cold-start evaluation pipeline; D3.2 primary result; D3.3 failure/stress-test map; D3.4 cross-archetype analysis.

Risk: operational data access. *Mitigation:* agreements initiated pre-award, letters accompany the application. If clinical data are delayed, H3a remains testable on open federal/cantonal and European surveillance series under the same rolling-origin restriction; §2.4 states the cost.

WP4: From predictive skill to operational value (M24–M48)

T4.1: Elicit operational losses and thresholds (M24–M32)

Structured **SHELF** elicitation ($n \approx 15\text{--}20$ across HUG emergency medicine, CASU-144 regulation and capacity management): elicit the consequences of early, late and unnecessary escalation, then derive thresholds from the losses. **The primary decision is fixed in advance:** trigger surge-capacity escalation when the forecast probability of the strained/critical state crosses the elicited threshold.

T4.2: Decision-analytic evaluation and equity audit (M30–M40)

Re-evaluate the WP3 forecasts with net-benefit/decision-curve analysis and value-of-information: do rankings change once consequences are incorporated? A model counts as useful only if its improvement crosses a decision-relevant threshold. Because operational records may encode structural differences, assess calibration and threshold performance across aggregate strata (age, sex, neighbourhood deprivation where appropriate): calibrated on average but miscalibrated for a relevant group is not operationally ready.

T4.3: Counterfactual analysis and prospective validation (M34–M48)

Estimate, for selected episodes, what would have changed had escalation followed the model's signal (a simple capacity model, propagated uncertainty: counterfactuals, not causal estimates). If authorised, the framework also runs in **shadow mode** (forecasts recorded, never used clinically); if not, the project is complete on retrospective evaluation and says so.

Deliverables. D4.1 elicited loss structure and equity audit; D4.2 decision-analytic evaluation; D4.3 counterfactual and prospective validation where feasible.

Methods, data protection and reproducibility. Version-controlled R/Python, a registered analysis plan for confirmatory comparisons; clinical data processed in the UNIGE secure environment under a new CCER approval with me as applicant; benchmark and software released openly, synthetic equivalents where possible.

Risks and fallback logic

The two central risks are scientific, and each yields a result rather than a stall: **evidence may be unusable** (WP1 establishes that boundary, WP3 quantifies the cost of ignoring it: a publishable negative result); **the state model may be weakly identifiable** (simulation first, ordinal fallback, the identifiability boundary reported as a result). None of the explicit fallbacks above converts an inconclusive analysis into a success claim.

Expected outputs. Approximately four to six papers and two durable open resources: the quantitative extraction benchmark and the evidence-to-prior reference framework. **I lead the methodological, benchmark and integrative outputs.**

2.3.3 Research environment, team and resources

Host: Institute of Global Health, Faculty of Medicine, University of Geneva, with a methodological collaboration with **Data Science for Digital Health** (Prof. Douglas Teodoro) and a

clinical collaboration with **HUG emergency medicine** (Prof. Thibaut Desmettre; Dr Robert Larribau, médecin responsable CASU-144).

The project needs three capabilities rarely housed together: quantitative methods, evidence synthesis, and access to a working emergency system. **Institute of Global Health:** epidemiology, infectious-disease modelling and automated evidence extraction in one institute; this is where the gap became visible. Independence is organisational: a programme led by me alongside the Institute's groups, not inside one (§2.6). **DS4DH** contributes the biomedical NLP WP1 needs; through GESICA I already work with the emergency-medicine partners, and the AI-in-EMS review [Edjinedja 2026] evidences the functioning collaboration.

Existing infrastructure. LiteRev-Evidence is a prototype (live federated search of the open literature, screening, living reviews, provenance-tracked extraction with quality scoring, pooling into parameter distributions, weather and surveillance connectors); WP1 and WP2 turn it into a validated instrument and test it. The legionellosis study (BASEC 2026-00324, ethics granted) supplies the contrasting archetype. **Computing:** UNIGE HPC (Baobab/Yggdrasil) and the University's secure environment for clinical data.

Commitments, by status.

Item	Status	Evidence / action
Legionellosis linked data	Secured	Ethics approval BASEC 2026-00324
Host institute (ISG)	Requested	Host confirmation, SNSF template
DS4DH methodological collaboration	Agreed	Collaboration letter requested
Institutional confirmation	Requested	General confirmation, UNIGE Vice-Rectorate
CASU-144 records (HUG-operated), primary outcome	Requested	Collaboration letter requested. Continuous, daily, near real-time, archive depth to confirm; HUG centrale access already established within GESICA (voice excluded); own agreement and CCER approval required
ED presentations, additional channel	Requested	Collaboration letter requested; daily historical availability to confirm
ICU occupancy data	Requested	Collaboration letter requested; fallback in WP3
Operational-data ethics	Requested	CCER submission, PI as applicant
Computing	Requested	UNIGE HPC access

The three operational-data rows are the key remaining feasibility items; §2.4 states their cost.

Resources requested. Project funds are capped at **CHF 250,000 over four years**; the applicant's salary is covered separately. Doctoral students and postdocs cannot be employed (2026 rules); support personnel can.

Item	Rationale
Scientific/technical collaborator, 50% over 48 months	WP1 benchmark extraction, WP3 harmonisation and evaluation pipeline, reproducibility
Second independent extractor	Contracted M3-M9; required for T1.2's dual-extraction design
Computing and data access	Evidence processing, Bayesian estimation, rolling-origin evaluation
Travel, conferences and incoming visits	Results presentation and sustained international collaboration
Other eligible direct costs	As justified in the final SNSF budget

Preparatory work before the grant starts. Advanced before month 1, outside grant funding: **data agreements** with HUG and 144/CASU; **CCER preparation with myself as applicant**; **infrastructure hardening** of LiteRev-Evidence before clinical data connect. Feasibility rests on one structural property (**the prototype already exists**) and on the fallbacks above, each leaving the central question answerable with the claim narrowed and the narrowing stated.

2.4 Schedule and milestones




Figure 2:

work packages, tasks and milestones over 48 months.

Who does what. I execute WP2 and WP3's confirmatory design personally, and lead WP1's protocol, WP3's analysis and WP4's decision work. One **scientific/technical collaborator** (support personnel, 50% over 48 months) carries three fully specified tasks: T1.2's second independent extractor, the T3.1-T3.2 harmonisation and rolling-origin pipeline, and the reproducibility engineering behind D1.1 and D3.1. No other personnel are requested.

The design avoids a serial chain: WP2 falls back to weakly informative priors if WP1 finds extraction inadequate, WP3 to open surveillance data if operational access is delayed, and WP4's decision analysis is retrospective. **What the data fallback costs:** Open surveillance series are not

emergency-system demand. Falling back **preserves the methodological test but narrows the outcome claim** to routinely observed crisis indicators, and turns most of WP4 illustrative. Hence the agreements are a pre-award action, not a risk to manage later.

Milestone	Month	Criterion
M1	9	Extraction benchmark released
M2	12	Regime/state representation passes identifiability criteria, or fallback selected
M3	20	Retrospective information set harmonised, or open-data fallback activated
M4	34	Primary cold-start hypothesis tested on the pre-specified ladder
M5	40	Decision relevance established
M6	48	Cross-domain and prospective validation where feasible

2.5 Relevance and impact

The project tests an assumption crisis forecasting makes informally: that accumulated evidence from elsewhere is useful when local outcome data are insufficient. The result is useful either way: a principled, auditable way to borrow if priors help; a stated boundary if not.

A null result changes identifiable practice. Modellers routinely seed early-crisis models with parameters lifted from a few studies, and evidence-synthesis platforms (including my own) assume pooled estimates are reusable as priors; a calibrated null tells those groups the practice buys nothing where it is most used.

Form of publication. Peer-reviewed articles (methods in biostatistics and evidence-synthesis venues, applied results in epidemiology and emergency/health-services journals), preprints on submission, open access throughout; the benchmark, prior library and evaluation code as **citable, versioned open datasets and software**; results returned to the operational partners.

Three durable resources remain, subordinate to the central question: **a quantitative extraction benchmark; an open reference framework** for evidence-informed latent-state forecasting; **a decision-analytic evaluation framework**. The project is not a new extraction system, an early-warning indicator or a forecasting competition.

Practical and societal impact. The contribution is preparedness, not a promised deployment. In a crisis's first weeks, emergency systems decide whether to open capacity, redistribute resources or escalate before local outcomes provide a reliable base; the project quantifies whether external evidence improves those decisions, and with how much residual uncertainty. Geneva anchors it; the contrasting domains show what transfers.

2.6 Relevance for personal career development

Why I am ready, and why now. I have the methods, the prototype and the domain access, used so far inside other people's programmes; what I have not had is a research agenda that is mine to set. My UNIGE trajectory is stated in full below: none of these programmes has this project's objective.

Programme	Objective	My role	Distinction from Ambizione
GESICA	Franco-Swiss crisis intelligence	Contribute or	Builds decision support; does not test evidence-derived priors
GeoAI4EI (Horizon Europe,	European epidemic-intelligence	Contribute	Builds infrastructure; does not test the inferential value of

Programme	Objective	My role	Distinction from Ambizione
2026-2030)	toolbox	or	borrowed evidence
Legionellosis study (ongoing)	Geneva waterborne outbreak determinants	Data/scientific lead	Supplies a contrasting validation extension
Ambizione	When evidence-derived priors help cold-start forecasting, and when they harm	PI, full-time	-

This project does not build another epidemic-intelligence platform. It uses existing infrastructures to test a question that is a work package of none of them.

The question and the methods are mine. It arose from a problem I met building LiteRev-Evidence: the platform could pool published estimates into priors, but nothing established whether those priors improved forecasting in a new setting. The methodological core came from fifteen years of quantitative finance, inherited from no group I have worked in; the platform is UNIGE intellectual property developed under my scientific direction.

Collaborators and their roles. Prof. Olivia Keiser leads the group in which I developed LiteRev, directs the MAS on which I teach, and leads the legionellosis study and GeoAI4EI; she co-authors my doctoral and post-doctoral output (the expected profile at this stage) and has **no scientific role in this project**. Prof. Douglas Teodoro leads GESICA and provides the WP1 biomedical-NLP collaboration.

The organisational position is the substance of the transition. Hosted by the **Institute of Global Health**, I will lead an independent research programme rather than continue as a member of an existing group, collaborating with the Institute's groups, including the one in which I developed LiteRev; the institutional guarantees of the PI's independence are documented in the required confirmation letters. The programme is executable by the PI and one budgeted collaborator: **no named individual is load-bearing, by design.**

Where this leads. By month 48 I intend to hold a validated answer, two open resources the field can build on, and a first-author methodological record in my own name: the basis for a group-leadership position or professorship. The follow-on programme fuses evidence streams of different velocity (literature, situational reporting, structured feeds) in one calibrated state model, conflict monitoring deciding how much each may say; this project supplies its evidence base. My research time is currently divided across four collaborative programmes; Ambizione consolidates it into one.

Bibliography

Status of this list

Two formal requirements from the call documents (Guidelines, "Bibliography"): (1) **Do not use "et al."**; give the full author list. The only exception is a publication involving a large international consortium with **over 50 authors**. (2) Give the full reference including **title, source and a DOI where possible**. Entries below marked [[FULL AUTHOR LIST REQUIRED]] must be completed from the publisher record; I have not invented names. Only three entries currently carry a DOI; add the rest.

Note also: **links to websites are not permitted** anywhere in the research plan or bibliography. The references below were **located and cross-checked against search results during drafting**. They are real works, and the attributions match what the sources report. They are **not yet verified against the articles themselves**; volume, issue, page numbers and exact author lists must be confirmed from the publisher record before submission, and three entries flagged **verify** authorship below are ones where the search returned the work but not a reliable author list. I have **not generated any citation I could not locate**. Where §2.1 makes a claim I could not source, it carries a [[...]] marker asking you to supply the reference from your own reading rather than a plausible-looking placeholder.

The same warning applies with more force to GESICA_Scientific_Base.md and ROADMAP.md in the LiteRev-Evidence repository: those are automatically generated and attribute specific performance figures to named papers. Verify every one before reuse.

LiteRev is the right instrument for completing this list; particularly for the novelty claim, where the novelty claim needs a documented systematic search you can report if challenged.

Surveillance and outbreak detection

- **[Farrington 1996]** Farrington CP, Andrews NJ, Beale AD, Catchpole MA. A statistical algorithm for the early detection of outbreaks of infectious disease. *J R Stat Soc A* 1996;159(3):547-563.
- **[Noufaily 2013]** Noufaily A, Enki DG, Farrington P, Garthwaite P, Andrews N, Charlett A. An improved algorithm for outbreak detection in multiple surveillance systems. *Stat Med* 2013;32(7):1206-1222. [online first 2012; check the year you cite]
- **[Le Strat 1999]** Le Strat Y, Carrat F. Monitoring epidemiologic surveillance data using hidden Markov models. *Stat Med* 1999;18(24):3463-3478.
- **[Watkins 2009]** Watkins RE, Eagleson S, Veenendaal B, Wright G, Plant AJ. Disease surveillance using a hidden Markov model. *BMC Med Inform Decis Mak* 2009;9:39. doi:10.1186/1472-6947-9-39

Prehospital data as a leading indicator

- **[Rosenkötter 2013]** Rosenkötter N, Ziemann A, Garcia-Castrillo Riesgo L, Gillet JB, Vergeiner G, Krafft T, Brand H. Validity and timeliness of syndromic influenza surveillance during the autumn/winter wave of A(H1N1) influenza 2009: results of emergency medical dispatch, ambulance and emergency department data from three European regions. *BMC Public Health* 2013;13:905. doi:10.1186/1471-2458-13-905 auditor-verified 30 Sep; confirm once against the publisher record
- **[EMS-ILI 2025]** [[AUTHOR LIST REQUIRED; from the publisher page; replace this citation key with the first author's name]]. Telephone calls to emergency medical service as a tool to predict influenza-like illness: a 10-year study. *Public Health* 2025;238:239-244. doi:10.1016/j.puhe.2024.12.021 final publication is 2025 (online Dec 2024); year corrected; Lombardy, seasons 2014–2024 excl. 2020–2022; note $R^2 \approx 0.50$ supports signal-not-equivalence
- **[Edjinedja 2026]** Edjinedja KL, Larribau R, **Orel E**, Cossus J, Elfahim O, Yendouname K, Vaussenat F, Teodoro D, Robert-Nicoud S, Barakat O, Desmettre T. Artificial intelligence in

emergency medical services for disasters and health emergencies: a systematic review.
[[submitted 2026; update status and venue before submission]]

Collaborative forecasting and its evaluation

- **[Cramer 2022]** Cramer EY, Ray EL, Lopez VK, et al. (*consortium >50 authors; shortening permitted by the guidelines; verify the count*). Evaluation of individual and ensemble probabilistic forecasts of COVID-19 mortality in the United States. *PNAS* 2022;119(15): e2113561119. doi:10.1073/pnas.2113561119
- **[Sherratt 2023]** Sherratt K, Gruson H, Grah R, et al. (*consortium >50 authors; shortening permitted; verify the count*). Predictive performance of multi-model ensemble forecasts of COVID-19 across European nations. *eLife* 2023. PMID 37083521.

Evidence as prior information

- **[Cook 2023]** Cook JD, [[FULL AUTHOR LIST REQUIRED]]. Bayesian forecasting of disease spread with little or no local data. *Scientific Reports* 2023;13. [[article number + doi; auditor-verified venue (s41598-023-35177-6); complete from the publisher record]]
 - The prior-art example §2.1.2 cites: Bayesian forecasting with informative priors under sparse local data. COLDSTART's gap is downstream of it; automated construction, propagated extraction/transportability uncertainty, strict historical information sets, harm detection.
- **[Ibrahim 2000]** Ibrahim JG, Chen M-H. Power prior distributions for regression models. *Statistical Science* 2000;15(1):46–60.
- **[Hobbs 2011]** Hobbs BP, Carlin BP, Mandrekar SJ, Sargent DJ. Hierarchical commensurate and power prior models for adaptive incorporation of historical information in clinical trials. *Biometrics* 2011;67(3):1047–1056. verify volume/pages
- **[Schmidli 2014]** Schmidli H, Gsteiger S, Roychoudhury S, O'Hagan A, Spiegelhalter D, Neuenschwander B. Robust meta-analytic-predictive priors in clinical trials with historical control information. *Biometrics* 2014;70(4):1023–1032.

Automated extraction

- **[Shankar 2026]** Shankar R, Lim A, Qian X. Performance of large language models in data extraction for evidence synthesis: a systematic review. *J Biomed Inform* 2026;181:105086. doi:10.1016/j.jbi.2026.105086 author list and details auditor-verified 30 Sep; confirm once against the publisher record
 - Key figures for §2.1.2: overall accuracy 47–99.9%; numerical items 47–88% vs categorical/string 74–96%; omissions 60–74% of errors; hallucination 0.08–6%.
- **[Orel 2023]** Orel E, Ciglenecki I, Thiabaud A, Temerev A, Calmy A, Keiser O, Merzouki A. An automated literature review tool (LiteRev) for streamlining and accelerating research using natural language processing and machine learning. *J Med Internet Res* 2023;25:e39736.

Transportability

- **[Bareinboim 2016]** Bareinboim E, Pearl J. Causal inference and the data-fusion problem. *PNAS* 2016;113(27):7345–7352. also: External validity: from do-calculus to transportability across populations, *Statistical Science* 2014
- **[Dahabreh 2019]** Dahabreh IJ, Petito LC, Robertson SE, Hernán MA, Steingrímsson JA. Towards causally interpretable meta-analysis: transporting inferences from multiple studies to a target population. arXiv:1903.11455. [verify final publication venue](#)
- **[Degtiar 2023]** Degtiar I, Rose S. A review of generalizability and transportability. *Annual Review of Statistics and Its Application* 2023;10:501–524. arXiv:2102.11904

Critical slowing down and resilience indicators

- **[Scheffer 2009]** Scheffer M, Bascompte J, Brock WA, Brovkin V, Carpenter SR, Dakos V, Held H, van Nes EH, Rietkerk M, Sugihara G. Early-warning signals for critical transitions. *Nature* 2009;461:53–59.
- **[O'Regan 2013]** O'Regan SM, Drake JM. Theory of early warning signals of disease emergence and leading indicators of elimination. *Theoretical Ecology* 2013;6:333–357. [verify pages](#)
- **[Brett 2018]** Brett TS, O'Dea EB, Marty É, Miller PB, Park AW, Drake JM, Rohani P. Anticipating epidemic transitions with imperfect data. *PLoS Comput Biol* 2018;14(6):e1006204.
- **[Southall 2021]** Southall E, Brett TS, Tildesley MJ, Dyson L. Early warning signals of infectious disease transitions: a review. *J R Soc Interface* 2021;18(182):20210555. **The review to work from:** it catalogues the candidate indicators and their known false-positive behaviour, which T2.4/T2.5 must confront rather than tune away.

Conformal prediction

- **[Angelopoulos 2023]** Angelopoulos AN, Bates S. Conformal prediction: a gentle introduction. *Foundations and Trends in Machine Learning* 2023;16(4):494–591. [verify volume/pages](#)
- **[Barber 2023]** Barber RF, Candès EJ, Ramdas A, Tibshirani RJ. Conformal prediction beyond exchangeability. *Annals of Statistics* 2023;51(2):816–845. key fixed (was [Xu 2023]); cited in T2.5

Regime switching and extremes

- **[Hamilton 1989]** Hamilton JD. A new approach to the economic analysis of nonstationary time series and the business cycle. *Econometrica* 1989;57(2):357–384.
- **[Coles 2001]** Coles S. *An Introduction to Statistical Modeling of Extreme Values*. Springer, 2001.
- **[Ranjbar 2022]** Ranjbar S, Cantoni E, Chavez-Demoulin V, Marra G, Radice R, Jatón-Ogay K. Modelling the extremes of seasonal viruses and hospital congestion: the example of flu in a Swiss hospital. *J R Stat Soc Ser C (Appl Stat)* 2022. doi:10.1111/rssc.12559. **The closest published precedent.** Discrete generalised Pareto model for extremes of flu-like hospital visits, three years of daily data from a large Swiss hospital (CHUV). Note the affiliations:

Cantoni is at the Research Center for Statistics, University of Geneva (the applicant's DEA supervisor); **Chavez-Demoulin is at UNIL**. Relevant as the nearest prior work to §2.1.3, not as a mobility lead; no research stay is proposed.

Entries added 30 Sep: cited in the text but previously MISSING from this list

Found by an external audit and verified: [Lazer 2014], [Winters 2018], [Lee 2021], [Höhle 2014] and [McGough 2020] were cited in §2.1/§2.3.2 with no bibliography entry. Author lists below are supplied where I am confident of them and flagged where I am not; verify all five against the publisher record like every other entry.

- **[Lazer 2014]** Lazer D, Kennedy R, King G, Vespignani A. The parable of Google Flu: traps in big data analysis. *Science* 2014;343(6176):1203–1205. doi:10.1126/science.1248506
- **[Höhle 2014]** Höhle M, an der Heiden M. Bayesian nowcasting during the STEC O104:H4 outbreak in Germany, 2011. *Biometrics* 2014;70(4):993–1002. doi:10.1111/biom.12194
verify pages
- **[McGough 2020]** McGough SF, Johansson MA, Lipsitch M, Menzies NA. Nowcasting by Bayesian smoothing: a flexible, generalizable model for real-time epidemic tracking. *PLoS Comput Biol* 2020;16(4):e1007735. doi:10.1371/journal.pcbi.1007735
- **[Winters 2018]** Winters BD, [[FULL AUTHOR LIST REQUIRED]]. Technological distractions (part 2): a summary of approaches to manage clinical alarms with intent to reduce alarm fatigue. *Crit Care Med* 2018;46(1):130–137. [[verify venue/volume; this entry sources the 72–99% false/non-actionable alarm figure in §2.1.4; confirm the figure appears in the final source]]
- **[Lee 2021]** [[FULL REFERENCE REQUIRED; the §2.1.4 claim that heat-warning thresholds calibrated to mortality diverge from morbidity/ED thresholds needs its actual source. If no "Lee 2021" exists in your reading, the auditor suggests Lung, Yeh & Hwang 2021 (PMC8471601) as a candidate that directly discusses mortality- vs morbidity-based heat thresholds; verify and rename the key accordingly]]

Evaluation

- **[Gneiting 2007]** Gneiting T, Raftery AE. Strictly proper scoring rules, prediction, and estimation. *J Am Stat Assoc* 2007;102(477):359–378.
- **[Vickers 2006]** Vickers AJ, Elkin EB. Decision curve analysis: a novel method for evaluating prediction models. *Med Decis Making* 2006;26(6):565–574.

Own work: cite from the CV, verify details

- **Orel E**, Esra R, Estill J, Thiabaud A, Marchand-Maillet S, Merzouki A, Keiser O. Prediction of HIV status based on socio-behavioural characteristics in East and Southern Africa. *PLoS ONE* 2022;17(3):e0264429.
- **Orel E**, Sobel J, Zanella MC, Iten A, Kaiser L, Keiser O, Vetter P. Difference in clinical severity between Delta variant and Omicron sub-lineages in a Swiss tertiary center. *CMI Communications* 2024.

- Estill J, Venkova-Marchevska P, Roelens M, **Orel E**, Temerev A, Flahault A, Keiser O. Future scenarios for the SARS-CoV-2 epidemic in Switzerland: an age-structured model. *F1000Research* 2020;9.
- Merzouki A, Estill J, **Orel E**, Tal K, Keiser O. Clusters of sub-Saharan African countries based on sociobehavioural characteristics and associated HIV incidence. *PeerJ* 2021;9:e10660.
- Esra R, Carstens J, Le Roux S, Mabuto T, Eisenstein M, Keiser O, **Orel E**, Merzouki A, De Voux L, Maskew M, Sharpey-Schafer K. Validation and improvement of a machine learning model to predict interruptions in antiretroviral treatment in South Africa. *J Acquir Immune Defic Syndr* 2023;92(1):42–49.
- Nwosu K, Fokam J, Wanda F, [[FULL AUTHOR LIST REQUIRED; no "et al." permitted]] (incl. **Orel E**). SARS-CoV-2 antibody seroprevalence and associated risk factors in an urban district in Cameroon. *Nat Commun* 2021;12:5851.
- Ng'ambi W, Estill J, Keiser O, Merzouki FA, **Orel E** (senior author). Machine learning-based classification of self-reported cardiovascular disease history in Africa using harmonised multi-country WHO STEPS surveys: 2014–2019. **Accepted for publication.** [[Venue, year, DOI once assigned; update on publication]]
- **[Schulte 2024]** Schulte F, Rösli M, Ragettli MS. Risk, attributable fraction and attributable number of cause-specific heat-related emergency hospital admissions in Switzerland. *Int J Public Health* 2024;69:1607349. doi:10.3389/ijph.2024.1607349
 - Swiss heat-ED evidence: weak/non-significant for respiratory causes overall; significant in the 75–84 subgroup. Grounds the §2.3.1 heat-arm caveat. *Sourced from the GESICA v5 inventory's reference list (ref 41); verify against the publisher record like every other entry.*
- **[Ragettli 2019]** Ragettli MS, Vicedo-Cabrera AM, Schindler C, Rösli M. Impact of the warm summer 2015 on emergency hospital admissions in Switzerland. *Environ Health* 2019;18:66. doi:10.1186/s12940-019-0507-1 (*GESICA v5 ref 42*)
- **[Zheng 2015]** Zheng XY, Ding H, Jiang LN, et al. [[consortium size < 50; full author list required]] Association between air pollutants and asthma emergency room visits and hospital admissions in time series studies: a systematic review and meta-analysis. *PLoS One* 2015;10(9):e0138146. doi:10.1371/journal.pone.0138146 (*GESICA v5 ref 44; pollution → asthma ED, international; the v5 report notes no recent Swiss study of acute episodes; a contribution space for this project*)

Still to source: supply from your own reading

- ~~Heat and emergency department presentations / EMS demand in Switzerland~~ **Closed:** Schulte 2024 and Ragettli 2019 above, from the applicant's own referenced inventory
- Out-of-hospital cardiac arrest and temperature (your platform's scenario documentation cites work here; verify each citation before reuse)
- ICU surge capacity modelling
- Legionellosis epidemiology and building water systems (SwissLEGIO; the CCER protocol bibliography is a starting point and is already checked)
- Swiss surveillance infrastructure and cantonal escalation procedures